

Anurag Gupta

Senior Software Engineer | Applied AI/LLM Systems

✉ anugupta0072@gmail.com ☎ 07496 981474 📍 London, UK 🌐 github.com/anurag1210

PROFESSIONAL SUMMARY

Senior Software Engineer with 17+ years of enterprise engineering experience, specialising in Agentic AI and LLM systems. Hands-on builder of production-grade agentic applications using LangGraph, LangChain, and the ReAct pattern — including multi-tool autonomous agents and RAG pipelines deployed with FastAPI and Docker. Deep understanding of agent orchestration, retrieval-augmented generation, NLP pipelines, prompt engineering, and AI evaluation methodologies. Strong foundation in regulated financial services environments (JPMorgan) with a proven ability to translate cutting-edge AI research into reliable, scalable production systems. Experienced collaborating with cross-functional stakeholders to translate complex requirements into delivered AI solutions.

PROFESSIONAL EXPERIENCE

Senior Software Engineer, JPMorgan Chase	06/2022 – Present
AI Project: Automated Regression Suite (AWS Bedrock · Claude Sonnet · Pydantic · Great Expectations · Databricks)	London, UK
• Conceived and independently built a production-grade AI-powered regression testing framework — the first of its kind within the team — autonomously validating data quality across 100+ production tables in a high-frequency deployment environment. • Engineered an LLM-driven metadata analysis pipeline where Claude Sonnet autonomously identifies business-critical columns per table — replacing hardcoded rules with context-aware, AI-driven column prioritisation. • Designed a dynamic validation code generation system where Claude Sonnet produces bespoke validation logic per table — eliminating manual test scripting entirely. • Implemented structured output via Pydantic to capture all LLM-generated outputs as validated JSON before downstream processing — ensuring production reliability. • Built end-to-end automated reporting pipeline ingesting Pydantic-structured results into Great Expectations, producing HTML reports consumed directly by Product Owners.	
Platform & Infrastructure:	
• Led modernisation of legacy on-prem data warehouse into a cloud-native Databricks analytics platform supporting large-scale workforce data across the firm. • Built scalable data pipelines using Python, PySpark, Databricks, and AWS (S3, Glue, Lambda, ECS) processing enterprise-scale analytical workloads. • Implemented prompt engineering strategies and output validation patterns to reduce hallucination risk in production financial systems. • Implemented Infrastructure as Code with Terraform and CI/CD automation (Spinnaker) to standardise deployments and reduce operational risk.	

Lead Software Engineer, JPMorgan Chase

04/2015 – 04/2022

London, UK

- Led design and delivery of enterprise MIS and firmwide BI reporting solutions supporting executive decision-making across multiple business lines.
- Built Python-based data pipelines consuming REST API endpoints, applying business logic transformations, and persisting to target databases — in a heavily regulated financial environment with strict audit and compliance requirements.
- Translated complex business requirements into robust data models and reporting pipelines; owned full development lifecycle from design through production deployment.
- Served as primary technical point of contact for business stakeholders, translating enterprise user requirements into scalable engineering solutions and communicating technical decisions to non-technical audiences.
- Mentored junior engineers and coordinated cross-functional delivery across data, technology, and business stakeholder teams.

EDUCATION

MSc Digital & Technology Solutions (Data Science & Machine Learning, L7),

2025 – Present

University of Exeter, UK

UK

- Coursework: Machine Learning, NLP, Model Evaluation, AI Systems Design

Bachelor of Engineering, Electrical & Electronics Engineering — 1st Class with Distinction, Visvesvaraya Technological University, Bangalore

Bangalore

TECHNICAL SKILLS

Agentic AI

LangGraph, ReAct Pattern, Multi-Agent Orchestration, Tool Use, Agent Memory & State

LLM & GenAI

LangChain, OpenAI API (GPT-4o, GPT-4o-mini), AWS Bedrock (Claude), Prompt Engineering, Guardrails, LLM-as-Judge Evaluation, Knowledge Graph Integration

Python

Production-grade Python, FastAPI, Async Python, Pydantic, PySpark, Pytest, scikit-learn

Other

SQL, PL/SQL, Unix Shell, REST API Design, Git, Jupyter

RAG & NLP

Retrieval-Augmented Generation, ChromaDB, OpenSearch, Embedding-based Semantic Search, Document Indexing, NLP Pipelines, Text Classification, Named Entity Recognition

Cloud & MLOps

Docker, AWS (S3, Glue, Lambda, EC2, ECS), Databricks, Terraform, CI/CD (Spinnaker), Streamlit, LangSmith

Data & ML

PyTorch (academic), Pandas, NumPy, TensorFlow, Supervised Learning, Feature Engineering, Model Evaluation

AI ENGINEERING PROJECTS

FinSight — AI-Powered Financial Research Platform,

LangChain · LangGraph · ReAct · ChromaDB · OpenAI GPT-4o · FastAPI · Docker · LangSmith

github.com/anurag1210/finsight Phase 1 — RAG Pipeline:

- Built complete ingestion pipeline: PDF extraction (pdfplumber), recursive character chunking (1000 chars / 200 overlap), OpenAI embedding generation, and ChromaDB vector storage.
- Key architectural discovery: switched from PyPDFLoader to pdfplumber after identifying character corruption — reinforcing a data quality over model quality principle critical in production AI systems.

- Implemented semantic search retrieval with metadata filtering, returning top-K relevant chunks with relevance scoring.
- Designed prompt engineering layer with XML-tagged context, financial analyst system prompt, and hallucination prevention guardrails.
- Implemented dual-layer security: input guards (prompt injection detection) and output guards (PII / system prompt leakage prevention).
- Built interactive Streamlit chat interface with session-based conversation history and real-time response display.
- Integrated LangSmith observability — every query automatically traced with latency, token usage, and cost per request with zero code instrumentation.

Phase 2 — Autonomous Agent:

- Evolved RAG into a fully autonomous agent using LangGraph with the ReAct pattern — enabling iterative plan, tool selection, observation, and self-correction loops.
- Designed multi-tool agentic framework with four specialised tools: SEC filing retrieval (RAG), financial data lookup, web search integration, and calculation engine — invoked autonomously based on query context.
- Implemented LangGraph state management for agent memory, multi-step reasoning loops, and conditional tool routing across complex financial queries.
- Developed automated evaluation pipeline across 12 test cases at 4 difficulty levels: keyword fact-matching (60.15%) and LLM-as-judge scoring (72.22%); achieved 100% accuracy on unanswerable questions.
- Containerised full stack with Docker and docker-compose; exposed RESTful API via FastAPI with Pydantic validation and API key authentication. Average query latency: 2.95 seconds end-to-end.

ConvFinQA — Conversational Financial QA Agent,

LangGraph · GPT-4o-mini · Tool Calling · Prompt Engineering · Python

github.com/anurag1210/convfinqa-agent

- Built a conversational agent for multi-hop numerical reasoning over financial documents using LangGraph with stateful conversation history — enabling Q2 to reference Q1 answers across chained questions.
- Implemented tool-augmented arithmetic via five calculator functions bound to GPT-4o-mini at temperature 0 — arithmetic never computed in the model's head.
- Achieved 70.8% accuracy across 185 questions on the ConvFinQA benchmark through three systematic prompt iterations — few-shot examples outperformed rule-based instructions (60.5%).
- Conducted systematic failure analysis identifying three distinct error categories: division order errors, percentage share vs change confusion, and conversational pronoun resolution failures.
- Implemented dual-metric evaluation: exact match for lookup questions and numerical accuracy within 1% tolerance for calculated answers.

CERTIFICATIONS

- AWS Certifications: AWS
Certified Cloud Practitioner |
AWS Certified Solutions
Architect